

Lab 6: Clustering

Name: **Krish Agarwal**

Reg No: **21112016**

Class: **4BSc DS A**

Objective :

1. Download the "Car Evaluation" Dataset from UCI Repository (<https://archive.ics.uci.edu/ml/datasets/Car+Evaluation> (<https://archive.ics.uci.edu/ml/datasets/Car+Evaluation>)). Remove the target 'Class Values' from the dataset while applying clustering algorithms.
2. Find the optimal number of clusters using Elbow and Silhouette Method.
3. Compare KMeans and Agglomerative Clustering methods for clustering the instances in the above dataset. Validate the optimal number of clusters found out in the previous question. **Hint:** Even if the algorithm does not require labels, for cross-checking of clustering values, you may use the labels.
4. Find what hyperparameters were suitable in KMeans (n_clusters, max_iter, init, algorithm)
5. Find what hyperparameters were suitable in Agglomerative Clustering (n_clusters, metric, linkage)
6. Plot Hierarchical Clustering (Dendrogram).
7. Compare the better clustering algorithm with any classification algorithm, and write your notes on the same.

Problem Definition :

Train 2 clustering models - **KMeans** and **Agglomerative** - for the given Car Evaluation dataset and compare both the models. After the comparison, validate the optimal number of clusters.

Observations :

- Via the elbow method, the optimal value of K is 4 as the graph bends at 4.
- Via the silhouette method, the optimal value of K is 4 as the maxima of the graph is at 4.
- The Silhouette Score of the KMeans model is 0.15756037323919975.
- The Silhouette Score of the Agglomerative model is 0.11147083427605417.
- It is noticed that the silhouette_scores are highest (0.172048) when there are 8 clusters, init = 'k-means++' and max iteration is either 100 or 200.
- It is observed that the silhouette score is the highest when the number of clusters are 2 and the linkage is average.
- The KMeans algorithm resulted in an optimal number of 8 clusters, while AGNES clustering found 4 clusters. KMeans partitions the data, while AGNES merges clusters hierarchically. The centroid distance metric used in KMeans is easier to comprehend, whereas reading dendrograms in AGNES requires experience and expertise. With KMeans, we had to specify the number of clusters, but in AGNES, there were no pre-determined cluster numbers.

References :

1. StackOverflow
2. GeekforGeeks
3. TutorialsPoint
4. Medium
5. W3School

Completion Status:

Question Number	Status
1	Completed
2	Completed
3	Completed
4	Completed
5	Completed
6	Completed
7	Completed

Code :

Q1) Download the "Car Evaluation" Dataset from UCI Repository. Remove the target 'Class Values' from the dataset while applying clustering algorithms.

In [34]:

```
1 # importing all the necessary libraries / modules
2 import pandas as pd
3 import numpy as np
4 import random
5 import warnings
6 warnings.filterwarnings("ignore")
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9 from krishKiLibrary import countUnique
10 from sklearn.preprocessing import OrdinalEncoder
11 from sklearn.cluster import KMeans
12 from sklearn.metrics import silhouette_score
13 from sklearn.cluster import AgglomerativeClustering
14 from scipy.cluster import hierarchy
15 from sklearn.decomposition import PCA
16 from scipy.cluster.hierarchy import linkage, dendrogram
17 from sklearn.metrics import silhouette_samples, silhouette_score
18 import matplotlib.cm as cm
19 from sklearn.model_selection import train_test_split
20 from sklearn.tree import DecisionTreeClassifier, plot_tree
21 from sklearn.metrics import accuracy_score
22 from sklearn.neighbors import KNeighborsClassifier
```

In [2]:

```
1 # --- Importing the dataframe ---
2 df = pd.read_csv("D:/Z/Downloads/car.data", names = ["Buying", "Maint", "Doors", "Pe
3
4 # --- Removing the target variable from the training dataset ---
5 df_training = df.drop('Class_Val', axis = 1)
```

In [3]:

```
1 df
```

Out[3]:

	Buying	Maint	Doors	Persons	Lug_boot	Safety	Class_Val
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc
...
1723	low	low	5more	more	med	med	good
1724	low	low	5more	more	med	high	vgood
1725	low	low	5more	more	big	low	unacc
1726	low	low	5more	more	big	med	good
1727	low	low	5more	more	big	high	vgood

1728 rows × 7 columns

In [4]:

```
1 # --- Counting the number of unique values in each column ---
2 for i in df_training.columns:
3     print(i, ":", countUnique(df_training, df_training[i].unique(), i))
```

Buying : {'vhigh': 432, 'high': 432, 'med': 432, 'low': 432}

Maint : {'vhigh': 432, 'high': 432, 'med': 432, 'low': 432}

Doors : {'2': 432, '3': 432, '4': 432, '5more': 432}

Persons : {'2': 576, '4': 576, 'more': 576}

Lug_boot : {'small': 576, 'med': 576, 'big': 576}

Safety : {'low': 576, 'med': 576, 'high': 576}

In [5]:

```
1 # --- Ordinally Encoding the values of the columns ---
2 buying_mapper = {'vhigh': 4, 'high': 3, 'med': 2, 'low': 1}
3 maint_mapper = {'vhigh': 4, 'high': 3, 'med': 2, 'low': 1}
4 doors_mapper = {2: 2, 3: 3, 4: 4, '5more': 5}
5 persons_mapper = {2: 1, 4: 2, 'more': 3}
6 lug_boot_mapper = {'big': 3, 'med': 2, 'small': 1}
7 safety_mapper = {'high': 3, 'med': 2, 'low': 1}
```

In [6]:

```
1 df_training["Buying"] = df_training["Buying"].replace(buying_mapper)
2 df_training["Maint"] = df_training["Maint"].replace(maint_mapper)
3 df_training["Doors"] = df_training["Doors"].replace(doors_mapper)
4 df_training["Persons"] = df_training["Persons"].replace(persons_mapper)
5 df_training["Lug_boot"] = df_training["Lug_boot"].replace(lug_boot_mapper)
6 df_training["Safety"] = df_training["Safety"].replace(safety_mapper)
```

In [7]:

```
1 df_training
```

Out[7]:

	Buying	Maint	Doors	Persons	Lug_boot	Safety
0	4	4	2	2	1	1
1	4	4	2	2	1	2
2	4	4	2	2	1	3
3	4	4	2	2	2	1
4	4	4	2	2	2	2
...
1723	1	1	5	3	2	2
1724	1	1	5	3	2	3
1725	1	1	5	3	3	1
1726	1	1	5	3	3	2
1727	1	1	5	3	3	3

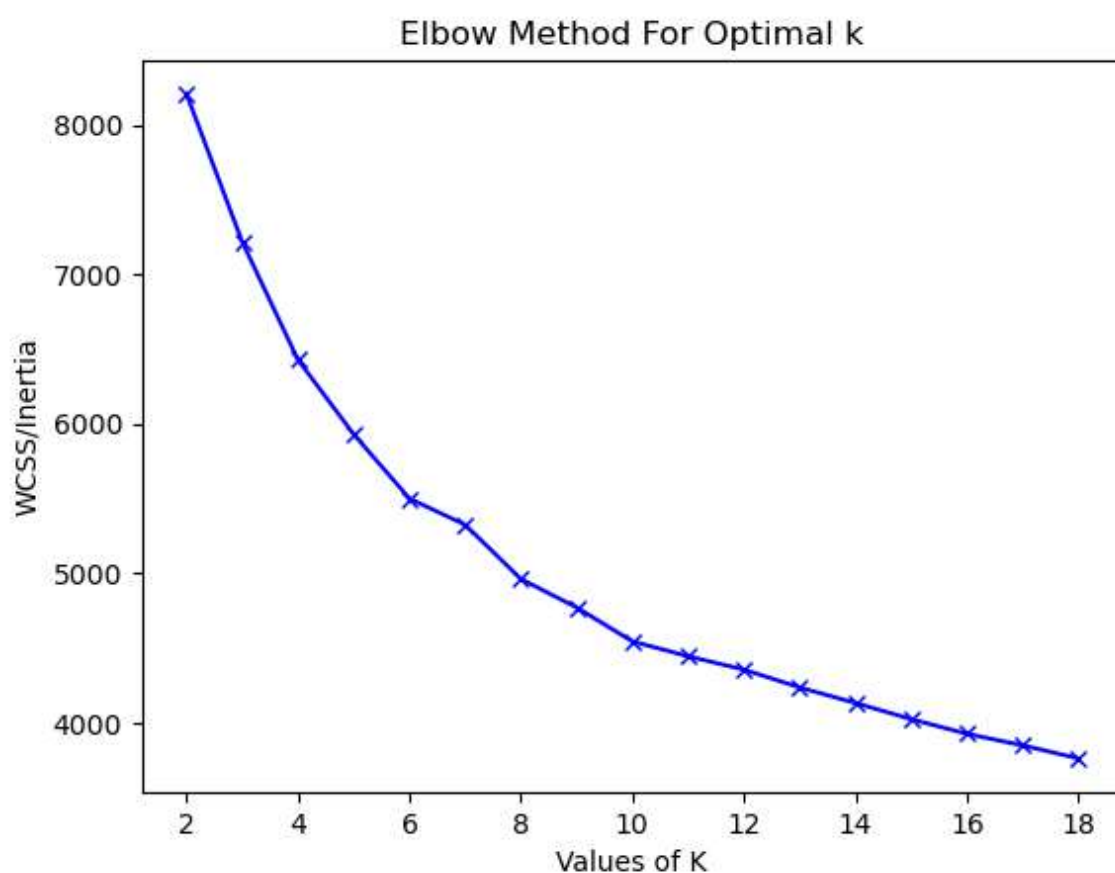
1728 rows × 6 columns

Q2) Find the optimal number of clusters using Elbow and Silhouette Method.

Elbow Method

In [8]:

```
1 # --- Plotting the elbow curve ---
2 wcss = []
3 K = range(2,19)
4
5 for num_clusters in K :
6     kmeans = KMeans(n_clusters = num_clusters, random_state = 42, n_init =1)
7     kmeans.fit(df_training)
8     wcss.append(kmeans.inertia_) # appending the wcss values into the list
9
10 plt.plot(K, wcss, 'bx-') # plotting the wcss values w.r.t. K
11 plt.xlabel('Values of K')
12 plt.ylabel('WCSS/Inertia')
13 plt.title('Elbow Method For Optimal k')
14 plt.show()
```

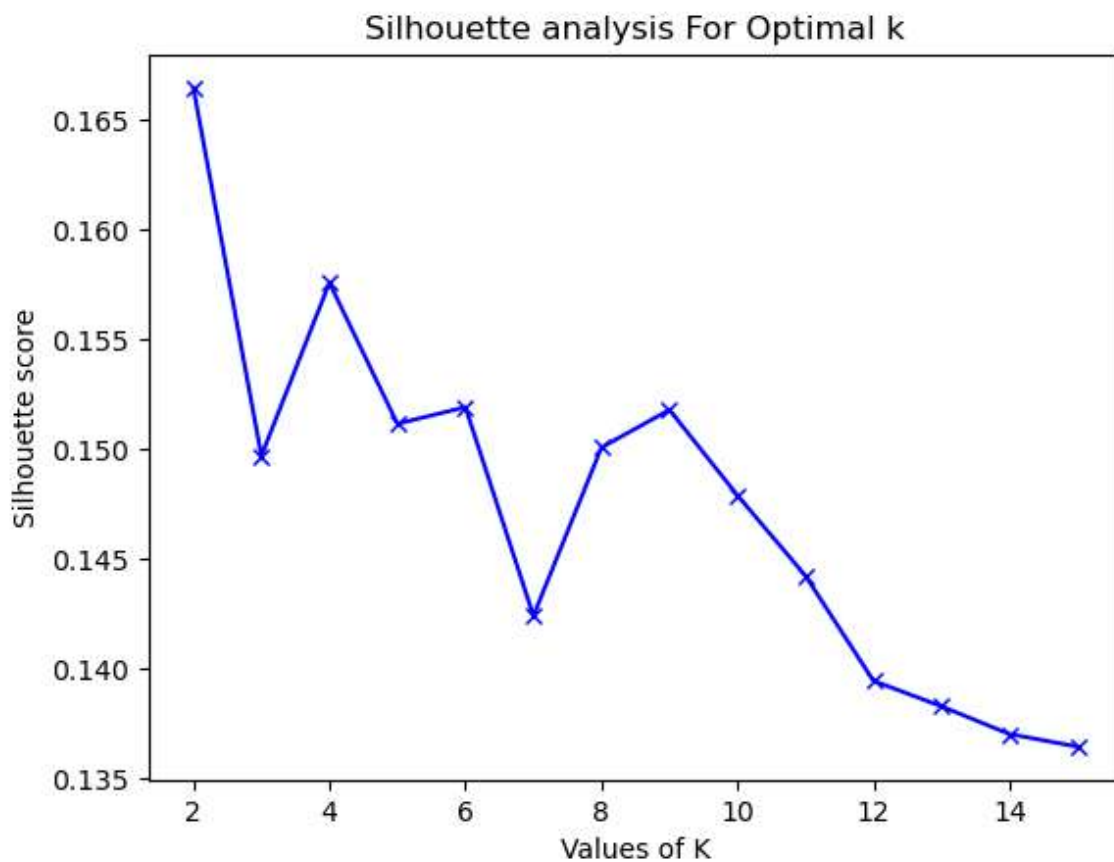


Inference: The optimal value of K is 4 as the graph bends there.

Silhouette Method

In [9]:

```
1 range_n_clusters = range(2, 16)
2 silhouette_avg = []
3
4 for num_clusters in range_n_clusters:
5     kmeans = KMeans(n_clusters = num_clusters, random_state = 42, n_init = 1) # init
6     kmeans.fit(df_training)
7     cluster_labels = kmeans.labels_
8     silhouette_avg.append(silhouette_score(df_training, cluster_labels))
9
10 plt.plot(range_n_clusters, silhouette_avg, 'bx-') # silhouette score
11 plt.xlabel('Values of K')
12 plt.ylabel('Silhouette score')
13 plt.title('Silhouette analysis For Optimal k')
14 plt.show()
```



Inference: The optimal value of K is 4 as the maxima of the graph is 4.

Q3) Agglomerative Clustering vs KMeans Clustering

KMeans Clustering

In [10]:

```
1 #Initialize the class object
2 kmeans = KMeans(n_clusters = 4).fit(df_training)
3 label = kmeans.predict(df_training)
4
5 # Calculating the aggregated silhouette score
6 silhouette_score_average = silhouette_score(df_training, label)
7 print('Silhouette Score (Accuracy): {}'.format(silhouette_score_average))
```

Silhouette Score (Accuracy): 0.1564952907335478

An Attempt to Plotting

In [11]:

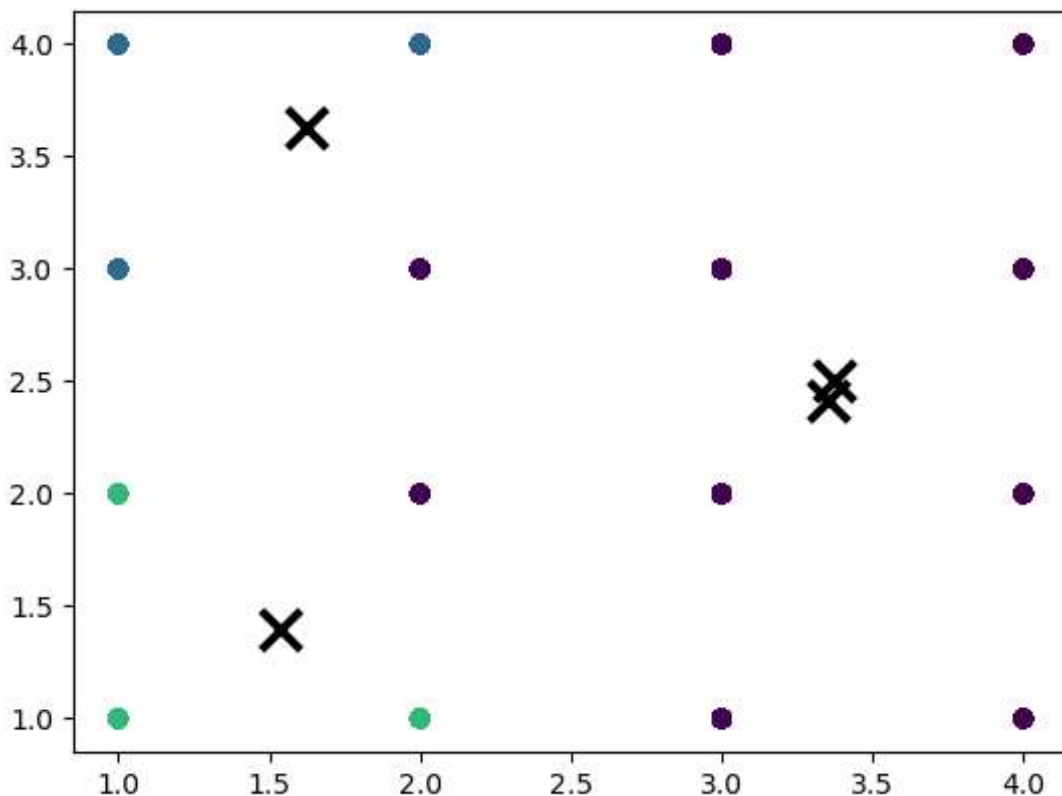
```
1 # Storing all the clusters in separate variables
2 label_0 = df_training[label == 0]
3 label_1 = df_training[label == 1]
4 label_2 = df_training[label == 2]
5 label_3 = df_training[label == 3]
```

In [12]:

```
1 # Plotting the clusters
2 plt.scatter(df_training.iloc[:, 0], df_training.iloc[:, 1], c = label, cmap = 'viridis')
3
4 # Plot the cluster centers as black dots
5 plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], marker = 'x')
```

Out[12]:

<matplotlib.collections.PathCollection at 0x2f228455fd0>



In [13]:

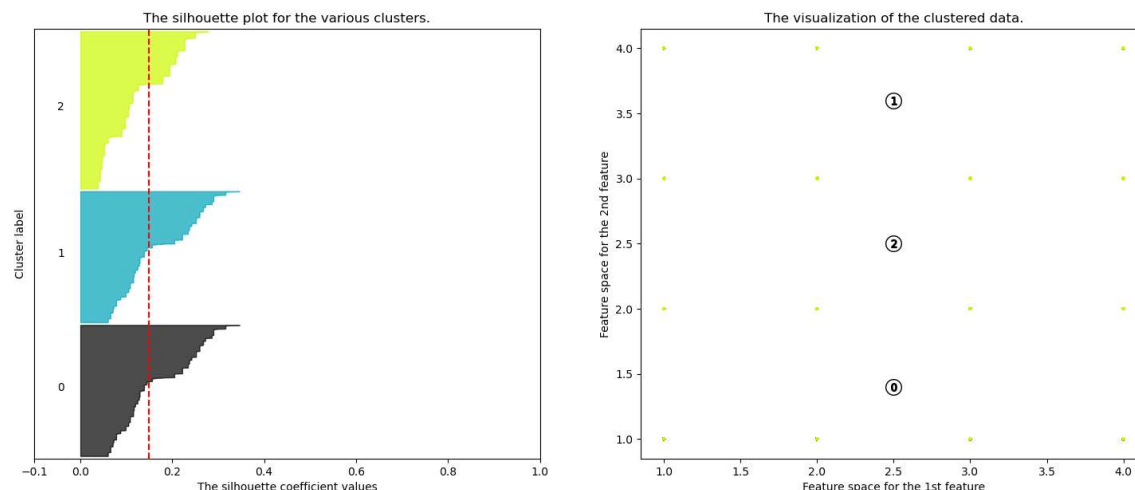
```
1 # Plotting KMEANS for various values of K
2 range_n_clusters = range(3,13)
3
4 for n_clusters in range_n_clusters:
5
6     # Create a subplot with 1 row and 2 columns
7     fig, (ax1, ax2) = plt.subplots(1, 2)
8     fig.set_size_inches(18, 7)
9
10    # The 1st subplot is the silhouette plot
11    # The silhouette coefficient can range from -1, 1 but in this example all
12    # lie within [-0.1, 1]
13    ax1.set_xlim([-0.1, 1])
14
15    # The (n_clusters+1)*10 is for inserting blank space between silhouette
16    # plots of individual clusters, to demarcate them clearly.
17    ax1.set_ylim([0, len(df_training) + (n_clusters + 1) * 10])
18
19    # Initialize the clusterer with n_clusters value and a random generator
20    # seed of 10 for reproducibility.
21    clusterer = KMeans(n_clusters=n_clusters, random_state=10)
22    cluster_labels = clusterer.fit_predict(df_training)
23
24    # The silhouette_score gives the average value for all the samples.
25    # This gives a perspective into the density and separation of the formed
26    # clusters
27    silhouette_avg = silhouette_score(df_training, cluster_labels)
28    print("For n_clusters =", n_clusters, "The average silhouette_score is :", silhc
29
30    # Compute the silhouette scores for each sample
31    sample_silhouette_values = silhouette_samples(df_training, cluster_labels)
32
33    y_lower = 10
34
35    for i in range(n_clusters):
36
37        # Aggregate the silhouette scores for samples belonging to
38        # cluster i, and sort them
39        ith_cluster_silhouette_values = \
40            sample_silhouette_values[cluster_labels == i]
41        ith_cluster_silhouette_values.sort()
42        size_cluster_i = ith_cluster_silhouette_values.shape[0]
43        y_upper = y_lower + size_cluster_i
44        color = cm.nipy_spectral(float(i) / n_clusters)
45        ax1.fill_betweenx(np.arange(y_lower, y_upper),
46            0, ith_cluster_silhouette_values,
47            facecolor=color, edgecolor=color, alpha=0.7)
48
49        # Label the silhouette plots with their cluster numbers at the middle
50        ax1.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))
51
52        # Compute the new y_lower for next plot
53        y_lower = y_upper + 10 # 10 for the 0 samples
54
55    ax1.set_title("The silhouette plot for the various clusters.")
56    ax1.set_xlabel("The silhouette coefficient values")
57    ax1.set_ylabel("Cluster label")
58
59    # The vertical line for average silhouette score of all the values
```

```

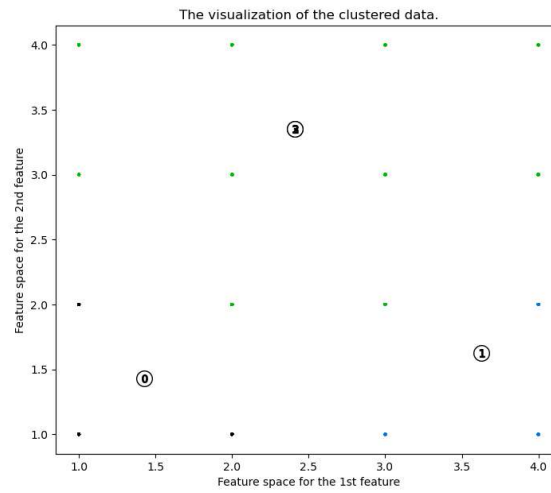
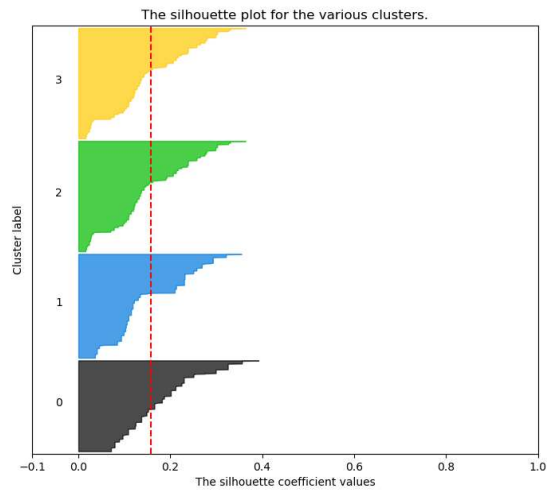
60 ax1.axvline(x=silhouette_avg, color="red", linestyle="--")
61 ax1.set_yticks([]) # Clear the yaxis labels / ticks
62 ax1.set_xticks([-0.1, 0, 0.2, 0.4, 0.6, 0.8, 1])
63
64 # 2nd Plot showing the actual clusters formed
65 colors = cm.nipy_spectral(cluster_labels.astype(float) / n_clusters)
66 ax2.scatter(df_training.iloc[:, 0], df_training.iloc[:, 1], marker='.', s=30, lw
67 c=colors, edgecolor='k')
68
69 # Labeling the clusters
70 centers = clusterer.cluster_centers_
71
72 # Draw white circles at cluster centers
73 ax2.scatter(centers[:, 0], centers[:, 1], marker='o',
74 c="white", alpha=1, s=200, edgecolor='k')
75
76 for i, c in enumerate(centers):
77     ax2.scatter(c[0], c[1], marker='.$d$' % i, alpha=1,
78 s=50, edgecolor='k')
79 ax2.set_title("The visualization of the clustered data.")
80 ax2.set_xlabel("Feature space for the 1st feature")
81 ax2.set_ylabel("Feature space for the 2nd feature")
82 plt.suptitle(("Silhouette analysis for KMeans clustering on sample data "
83 "with n_clusters = %d" % n_clusters),
84 fontsize=14, fontweight='bold')
85
86 For n_clusters = 3 The average silhouette_score is : 0.1496903095639953
87 For n_clusters = 4 The average silhouette_score is : 0.15756037323919983
88 For n_clusters = 5 The average silhouette_score is : 0.15168620164271934
89 For n_clusters = 6 The average silhouette_score is : 0.16154197589141794
90 For n_clusters = 7 The average silhouette_score is : 0.15929390277317174
91 For n_clusters = 8 The average silhouette_score is : 0.17204794782267452
92 For n_clusters = 9 The average silhouette_score is : 0.1582127116740343
93 For n_clusters = 10 The average silhouette_score is : 0.14964182627756134
94 For n_clusters = 11 The average silhouette_score is : 0.14260146059974588
95 For n_clusters = 12 The average silhouette_score is : 0.14131637907315775

```

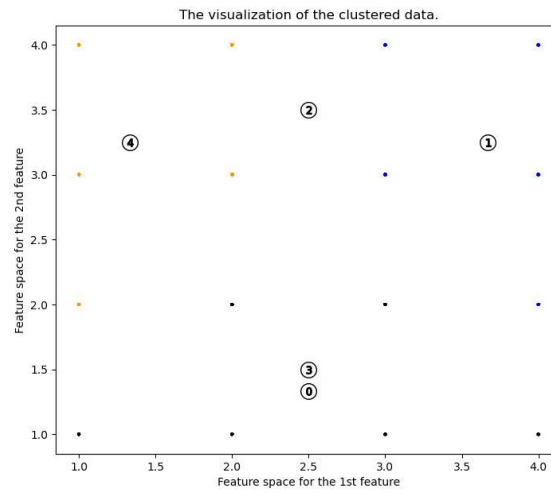
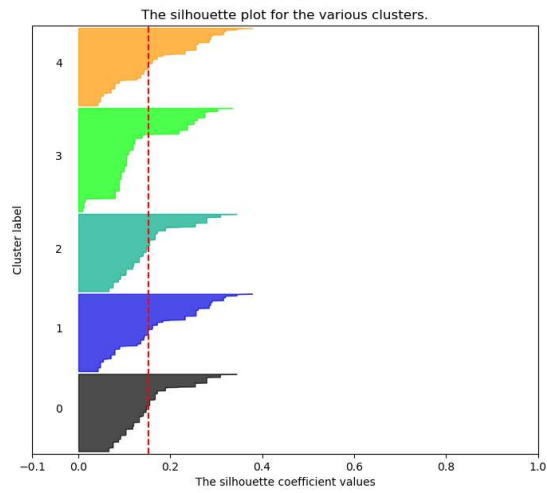
Silhouette analysis for KMeans clustering on sample data with n_clusters = 3



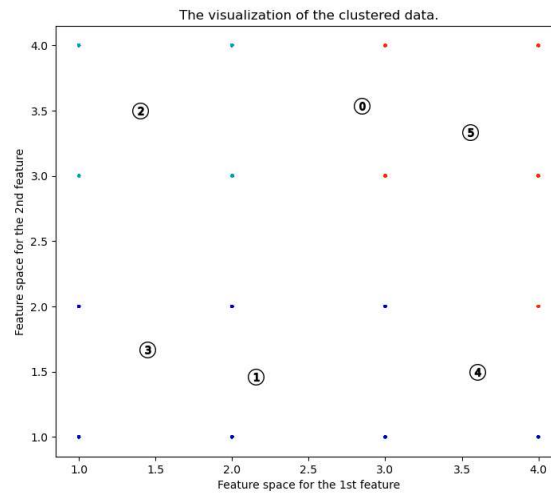
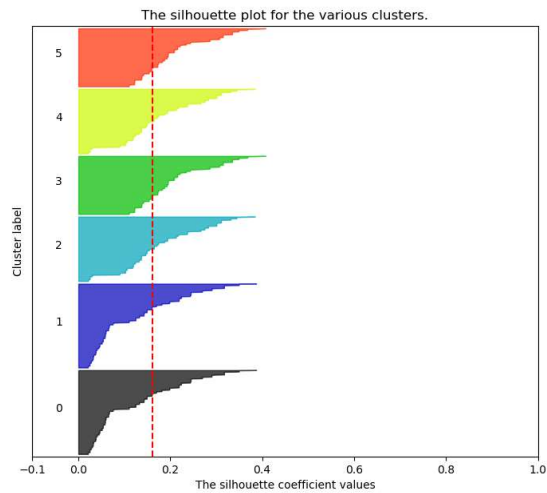
Silhouette analysis for KMeans clustering on sample data with $n_clusters = 4$



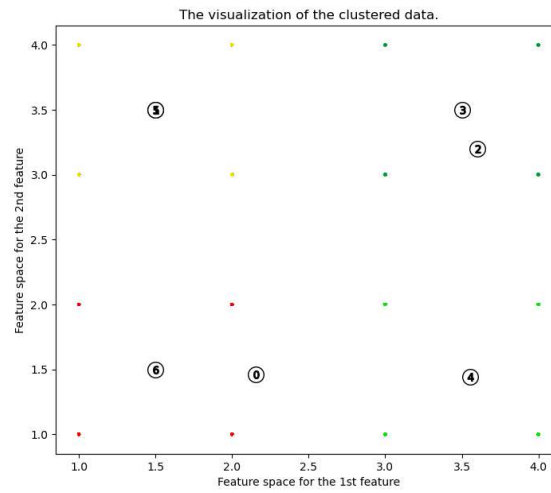
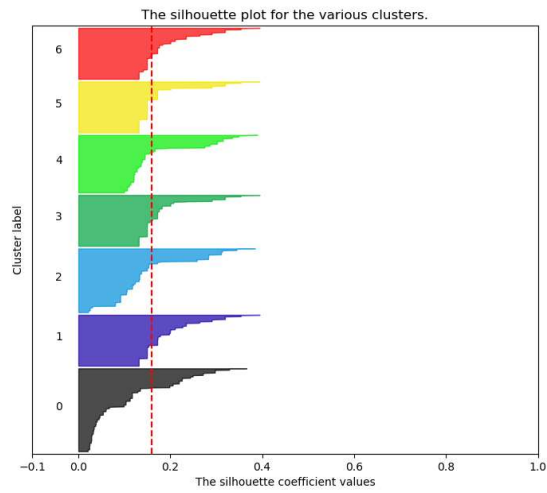
Silhouette analysis for KMeans clustering on sample data with $n_clusters = 5$



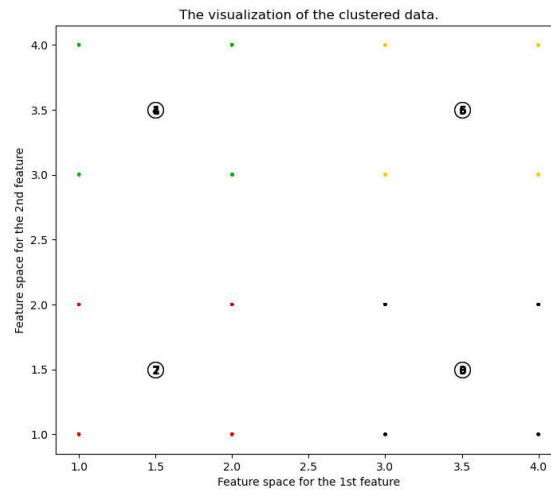
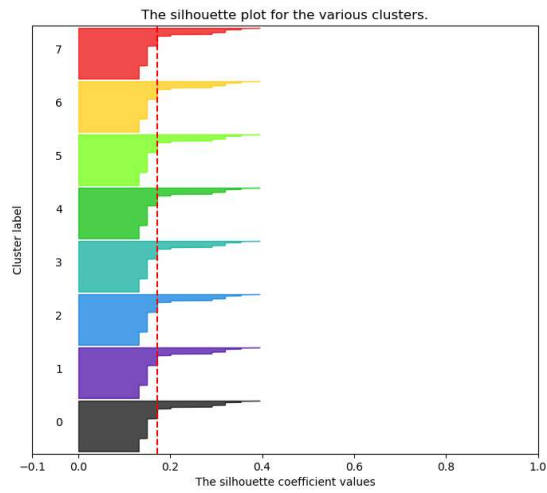
Silhouette analysis for KMeans clustering on sample data with $n_clusters = 6$



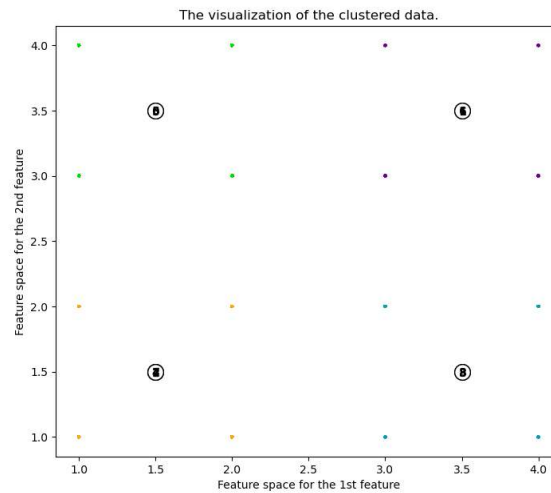
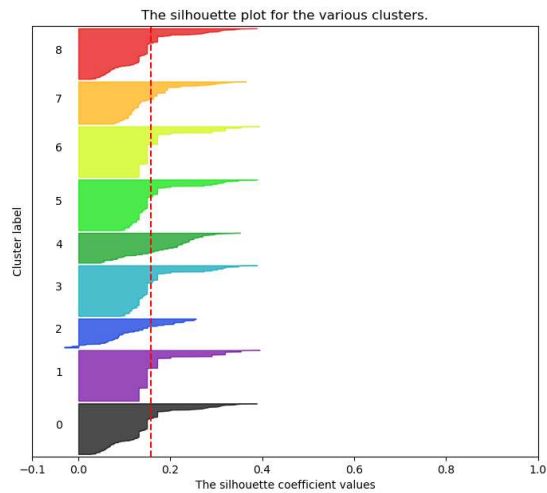
Silhouette analysis for KMeans clustering on sample data with $n_clusters = 7$



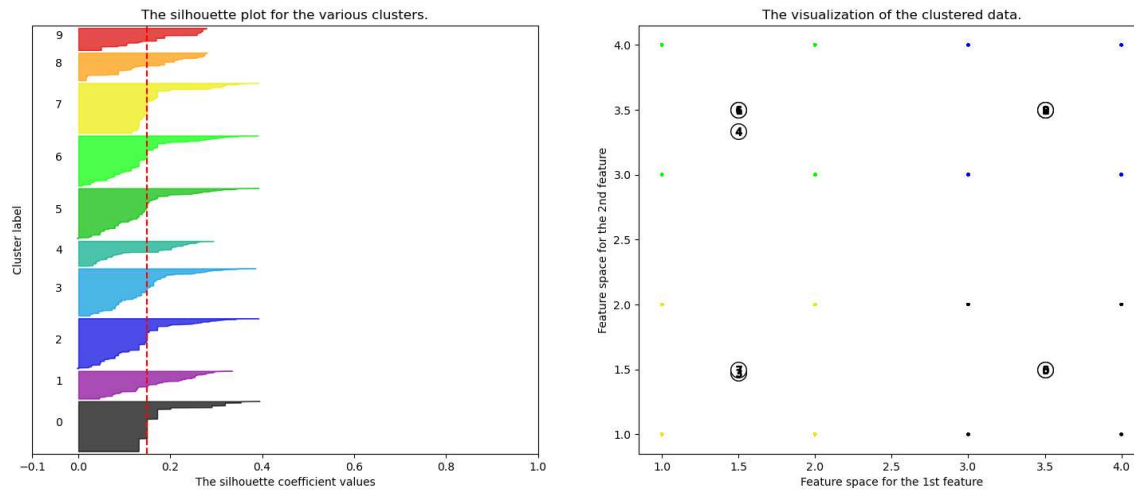
Silhouette analysis for KMeans clustering on sample data with $n_clusters = 8$



Silhouette analysis for KMeans clustering on sample data with $n_clusters = 9$

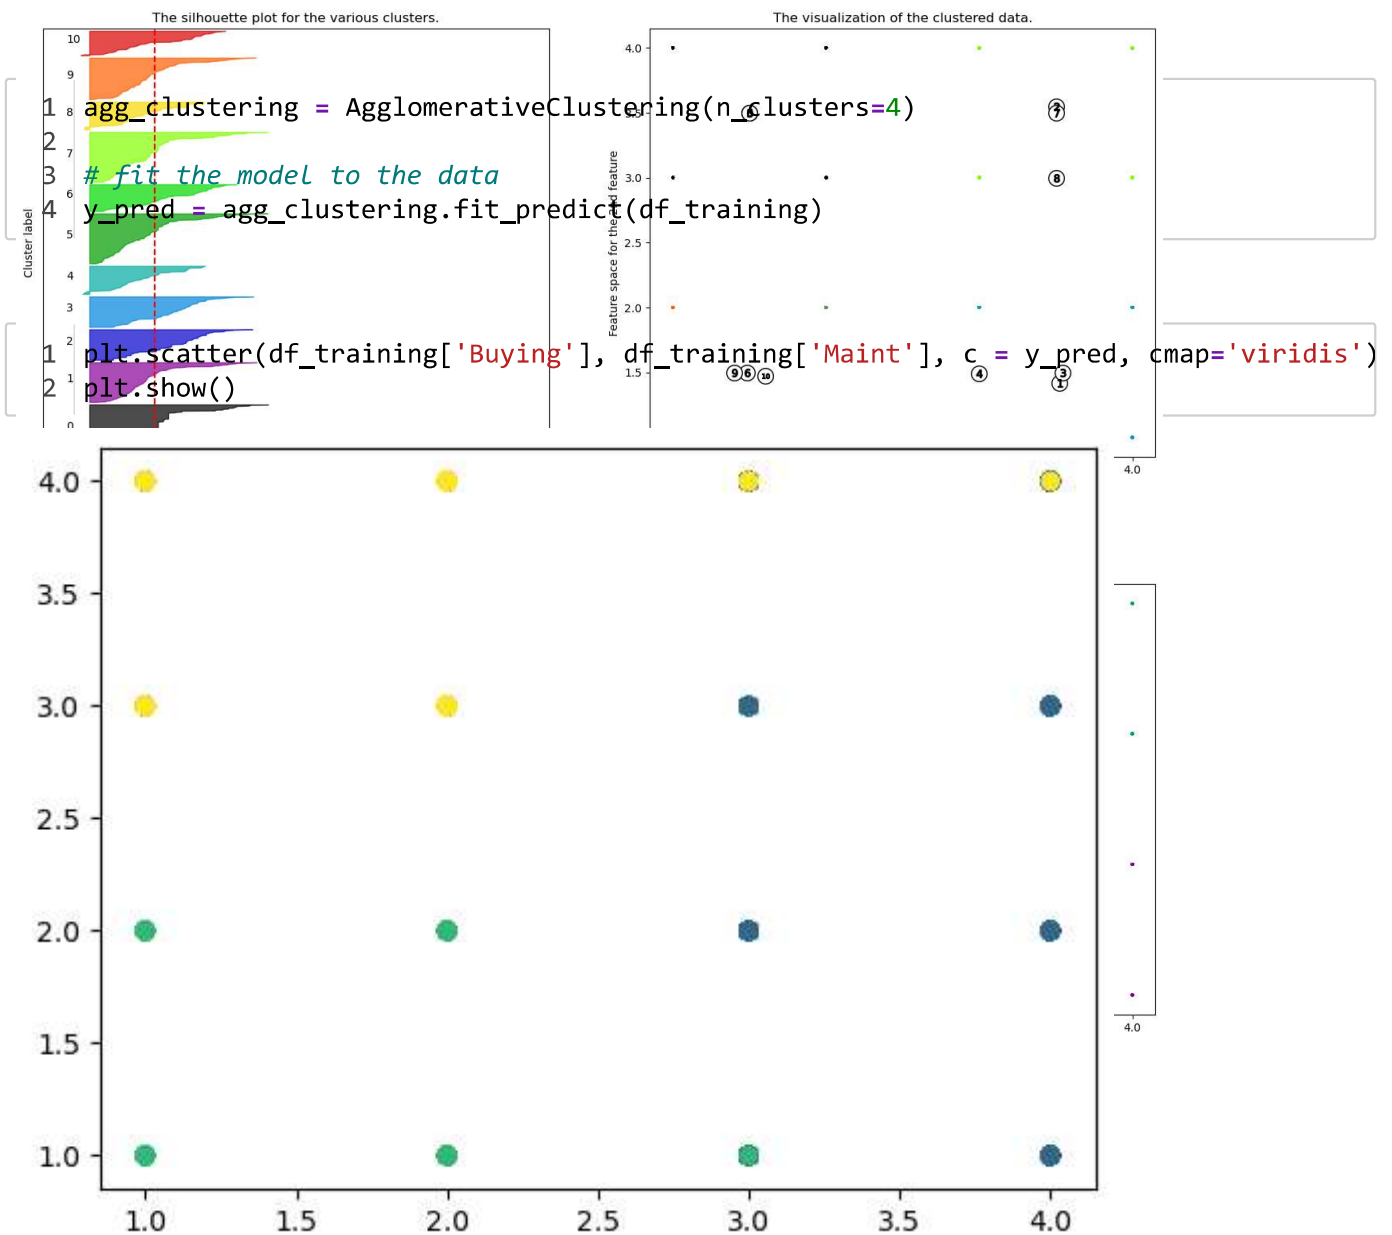


Silhouette analysis for KMeans clustering on sample data with n_clusters = 10



Agglomerative Clustering

Silhouette analysis for KMeans clustering on sample data with n_clusters = 11



In [16]:

```
1 print("Silhouette Score (Accuracy): ", silhouette_score(df_training, y_pred))
```

Silhouette Score (Accuracy): 0.11147083427605417

Q4) Suitability of Hyperparameters in KMeans Clustering

In [17]:

```
1 # Setting up hyperparameters
2 n_clusters_ = range(2, 16)
3 init_ = ['k-means++', 'random']
4 max_iter_ = [100, 200, 300, 400, 500]
5 random_state_ = [234, 325, 98, 235, 324, 3]
```

In [18]:

```
1 # Creating a dataframe to store all the possible combinations
2 df_ = pd.DataFrame(columns = ['Algorithm', 'init', 'n_clusters', 'max_iter', 'random
3
4 # Calculating all the possible silhouette scores for the above set hyper parameters
5 for i in init_:
6     for j in max_iter_:
7         for k in random_state_:
8             for l in n_clusters_:
9                 kmeans = KMeans(init = i, max_iter = j, random_state = k, n_clusters
10                 pred = kmeans.predict(df_training)
11                 ss = silhouette_score(df_training, pred)
12
13                 dict_ = {}
14                 dict_['Algorithm'] = "KMeans"
15                 dict_['n_clusters'] = l
16                 dict_['init'] = i
17                 dict_['random_state'] = k
18                 dict_['max_iter'] = j
19                 dict_['silhouette_score'] = ss
20
21                 df_ = df_.append(dict_, ignore_index = True)
```

In [19]:

```
1 df_
```

Out[19]:

	Algorithm	init	n_clusters	max_iter	random_state	silhouette_score
0	KMeans	k-means++	2	100	234	0.166422
1	KMeans	k-means++	3	100	234	0.149690
2	KMeans	k-means++	4	100	234	0.161509
3	KMeans	k-means++	5	100	234	0.151145
4	KMeans	k-means++	6	100	234	0.161870
...
835	KMeans	random	11	500	3	0.141286
836	KMeans	random	12	500	3	0.136150
837	KMeans	random	13	500	3	0.138239
838	KMeans	random	14	500	3	0.133412
839	KMeans	random	15	500	3	0.137375

840 rows × 6 columns

In [20]:

```
1 # Following are 10 instances with the largest silhouette scores
2 df_.iloc[df_['silhouette_score'].nlargest(10).index]
```

Out[20]:

	Algorithm	init	n_clusters	max_iter	random_state	silhouette_score
6	KMeans	k-means++	8	100	234	0.172048
34	KMeans	k-means++	8	100	98	0.172048
48	KMeans	k-means++	8	100	235	0.172048
62	KMeans	k-means++	8	100	324	0.172048
76	KMeans	k-means++	8	100	3	0.172048
90	KMeans	k-means++	8	200	234	0.172048
118	KMeans	k-means++	8	200	98	0.172048
132	KMeans	k-means++	8	200	235	0.172048
146	KMeans	k-means++	8	200	324	0.172048
160	KMeans	k-means++	8	200	3	0.172048

Inference: It is noticed that the silhouette_scores are highest (0.172048) when there are 8 clusters, init = 'k-means++' and max iteration is either 100 or 200.

Q5) Suitability of Hyperparameters in Agglomerative Clustering

In [21]:

```
1 # Setting up hyperparameters
2 n_clusters_ = range(2, 16)
3 linkage_ = ['ward', 'ward', 'complete', 'average', 'single']
4 random_state_ = [234, 325, 98, 235, 324, 3]
```

In [22]:

```
1 # Creating a dataframe to store all the possible combinations
2 df__ = pd.DataFrame(columns = ['Algorithm', 'linkage', 'n_clusters', 'silhouette_score'])
3
4 # Calculating all the possible silhouette scores for the above set hyper parameters
5 for i in linkage_:
6     for l in n_clusters_:
7         agg_clustering = AgglomerativeClustering(n_clusters = l, linkage = i)
8         pred = agg_clustering.fit_predict(df_training)
9         ss = silhouette_score(df_training, pred)
10
11         dict_ = {}
12         dict_['Algorithm'] = "Agglomerative Clustering"
13         dict_['n_clusters'] = l
14         dict_['linkage'] = i
15         dict_['silhouette_score'] = ss
16
17         df__ = df__.append(dict_, ignore_index = True)
```

In [23]:

```
1 df__
```

Out[23]:

	Algorithm	linkage	n_clusters	silhouette_score
0	Agglomerative Clustering	ward	2	0.116495
1	Agglomerative Clustering	ward	3	0.115315
2	Agglomerative Clustering	ward	4	0.111471
3	Agglomerative Clustering	ward	5	0.110009
4	Agglomerative Clustering	ward	6	0.104988
...
65	Agglomerative Clustering	single	11	-0.043878
66	Agglomerative Clustering	single	12	-0.045161
67	Agglomerative Clustering	single	13	-0.046390
68	Agglomerative Clustering	single	14	-0.047671
69	Agglomerative Clustering	single	15	-0.049013

70 rows × 4 columns

In [24]:

```
1 # Following are 10 instances with the largest silhouette scores
2 df__.iloc[df__['silhouette_score'].nlargest(10).index]
```

Out[24]:

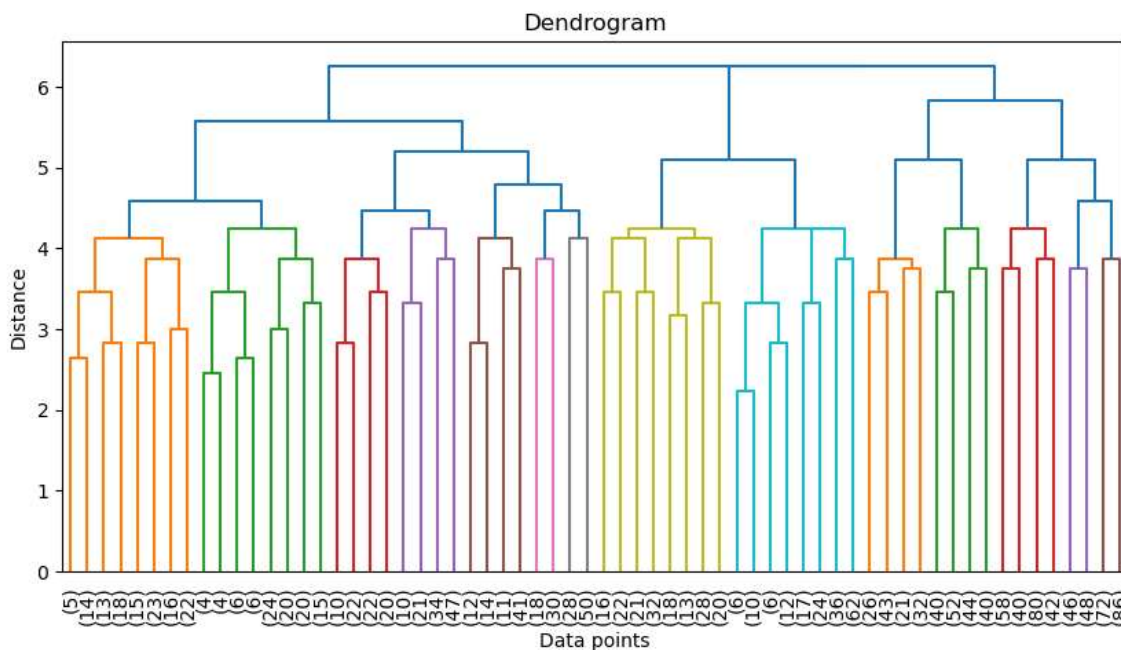
	Algorithm	linkage	n_clusters	silhouette_score
42	Agglomerative Clustering	average	2	0.150735
0	Agglomerative Clustering	ward	2	0.116495
14	Agglomerative Clustering	ward	2	0.116495
1	Agglomerative Clustering	ward	3	0.115315
15	Agglomerative Clustering	ward	3	0.115315
2	Agglomerative Clustering	ward	4	0.111471
16	Agglomerative Clustering	ward	4	0.111471
3	Agglomerative Clustering	ward	5	0.110009
17	Agglomerative Clustering	ward	5	0.110009
43	Agglomerative Clustering	average	3	0.109328

Inference: It is observed that the silhouette score is the highest when the number of clusters are 2 and the linkage is average.

Q6) Hierarchical Clustering (Dendrogram)

In [25]:

```
1 link = linkage(df_training, method = "complete")
2
3 # Plot dendrogram
4 plt.figure(figsize = (10, 5))
5 dendrogram(link, leaf_font_size=10, truncate_mode = "level", p = 5)
6 plt.xlabel("Data points")
7 plt.ylabel("Distance")
8 plt.title("Dendrogram")
9 plt.show()
```



Q7) Compare the better clustering algorithm with any classification algorithm, and write your notes on the same.

DecisionTreeClassifier

In [28]:

```
1 X = df_training.copy()
2 y = df['Class_Val']
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_st
```

In [29]:

```
1 # Build the decision tree
2 tree = DecisionTreeClassifier(max_depth = 5, min_samples_split = 5)
3
4 # Train the decision tree
5 tree.fit(X_train, y_train)
6
7 # Make predictions on the test set
8 y_pred = tree.predict(X_test)
```

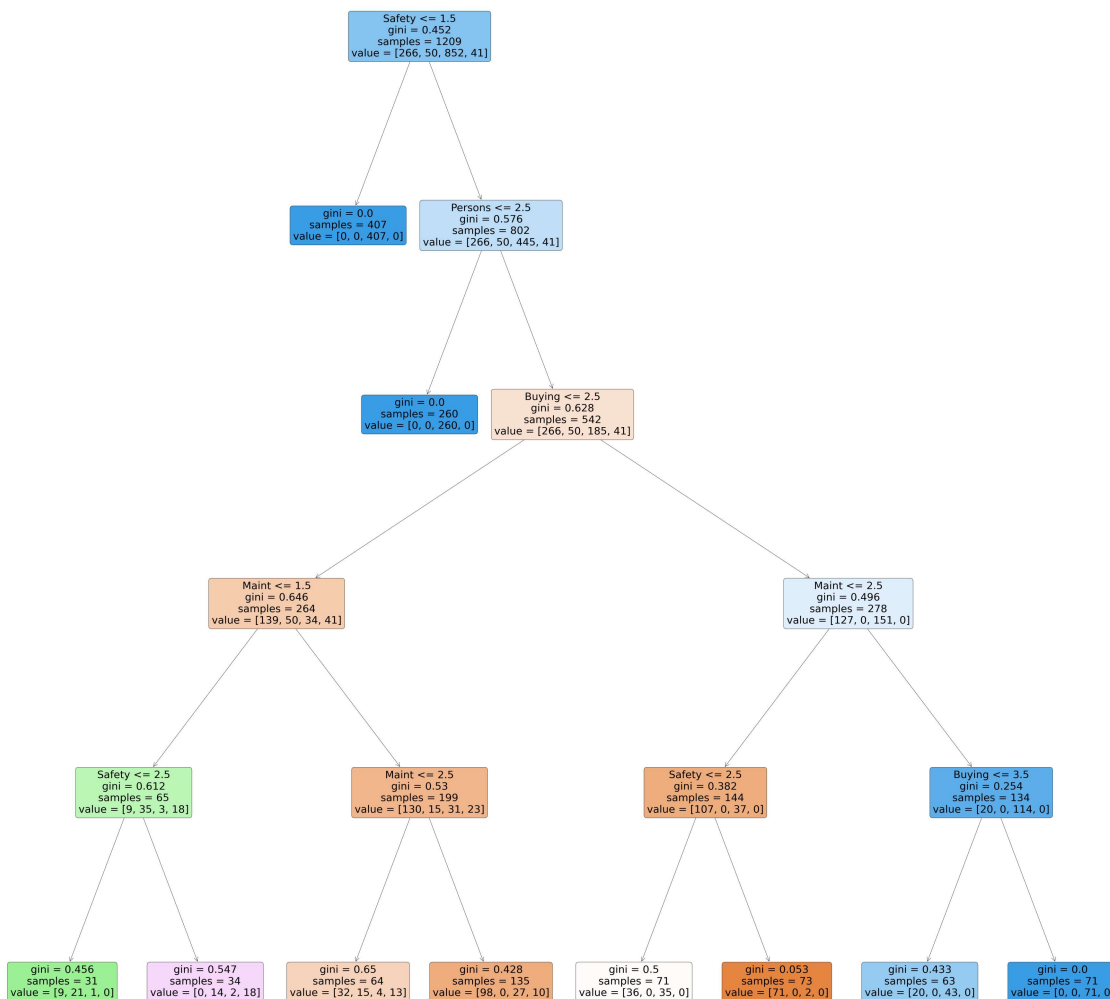
In [31]:

```
1 # Evaluate the accuracy of the model
2 accuracy = accuracy_score(y_test, y_pred)
3 print("Accuracy of Decision Tree:", accuracy)
```

Accuracy of Decision Tree: 0.8458574181117534

In [33]:

```
1 # Plotting the decision tree
2 fig, ax = plt.subplots(figsize=(80, 80))
3 plot_tree(tree, filled=True, rounded=True, ax=ax, feature_names=X.columns)
4 plt.show()
```



KNN Classifier

In [35]:

```
1 # Train the KNN classifier
2 knn = KNeighborsClassifier(n_neighbors=5)
3 knn.fit(X_train, y_train)
4
5 # Make predictions on the test set
6 y_pred = knn.predict(X_test)
7
8 # Evaluate the accuracy of the model
9 accuracy = accuracy_score(y_test, y_pred)
10 print("Accuracy of KNN:", accuracy)
```

Accuracy of KNN: 0.9402697495183044

Inference: The KMeans algorithm resulted in an optimal number of 8 clusters, while AGNES clustering found 4 clusters. KMeans partitions the data, while AGNES merges clusters hierarchically. The centroid distance metric used in KMeans is easier to comprehend, whereas reading dendrograms in AGNES requires experience and expertise. With KMeans, we had to specify the number of clusters, but in AGNES, there were no pre-determined cluster numbers.
